

Project 2: Data Mining, Classification, Prediction

SDS322E

Mining, Classification, Prediction

Lindsay Chu

Introduction

My dataset, “credit_card_data”, comes from Kaggle; it summarizes credit card usage of 8950 unique users based on 18 variables. Most are fairly straightforward, such as CUSTID, BALANCE, PURCHASES, and CREDIT LIMIT; others are scores from 0-1 and measure frequency, such as BALANCE_FREQUENCY (how often the balance is updated) and CASH_ADVANCE_FREQUENCY (how often payments are made in cash in advance).

I created the categorical variable, “AT_RISK”, which assigns the value “TRUE” to customers whose credit utilization rates are greater than 30%, and “FALSE” otherwise. Credit utilization was calculated by dividing “BALANCE” by “CREDIT LIMIT”, and the minimum threshold of 30% was chosen based on external research on “good” versus “poor” credit scores. Consequently, there are 2347 users “at-risk” of low credit scores / default (“TRUE” category), and 2128 users who are not at risk.

NOTE: I ended up cutting dataset in half (i.e. using the first 4475 rows) because R would not load any visualizations with all 8950 observations.

```
library(tidyverse)
library(dplyr)

# read dataset
ccdata <- read_csv("credit_card_data.csv")

glimpse(ccdata)

## Rows: 8,950
## Columns: 18
## $ CUST_ID          <chr> "C10001", "C10002", "C10003", "C10004~
## $ BALANCE          <dbl> 40.90075, 3202.46742, 2495.14886, 166~
## $ BALANCE_FREQUENCY <dbl> 0.818182, 0.909091, 1.000000, 0.63636~
## $ PURCHASES        <dbl> 95.40, 0.00, 773.17, 1499.00, 16.00, ~
## $ ONEOFF_PURCHASES <dbl> 0.00, 0.00, 773.17, 1499.00, 16.00, 0~
## $ INSTALLMENTS_PURCHASES <dbl> 95.40, 0.00, 0.00, 0.00, 0.00, 1333.2~
## $ CASH_ADVANCE      <dbl> 0.0000, 6442.9455, 0.0000, 205.7880, ~
## $ PURCHASES_FREQUENCY <dbl> 0.166667, 0.000000, 1.000000, 0.08333~
## $ ONEOFF_PURCHASES_FREQUENCY <dbl> 0.000000, 0.000000, 1.000000, 0.08333~
## $ PURCHASES_INSTALLMENTS_FREQUENCY <dbl> 0.083333, 0.000000, 0.000000, 0.00000~
## $ CASH_ADVANCE_FREQUENCY <dbl> 0.000000, 0.250000, 0.000000, 0.08333~
## $ CASH_ADVANCE_TRX   <dbl> 0, 4, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0~
## $ PURCHASES_TRX      <dbl> 2, 0, 12, 1, 1, 8, 64, 12, 5, 3, 12, ~
## $ CREDIT_LIMIT       <dbl> 1000, 7000, 7500, 7500, 1200, 1800, 1~
## $ PAYMENTS           <dbl> 201.8021, 4103.0326, 622.0667, 0.0000~
```

```
## $ MINIMUM_PAYMENTS          <dbl> 139.50979, 1072.34022, 627.28479, NA,~
## $ PRC_FULL_PAYMENT          <dbl> 0.000000, 0.222222, 0.000000, 0.00000~
## $ TENURE                    <dbl> 12, 12, 12, 12, 12, 12, 12, 12, 12, 1~
```

```
# tidying data
```

```
# select first half of the dataset because R would not load any visualizations with all of the observat
```

```
ccdata <- ccdata[1:4475,] %>%
```

```
# shorten variable names for readability in visualizations
```

```
rename(BAL_FREQ = BALANCE_FREQUENCY,
       ONEOFF = ONEOFF_PURCHASES,
       INSTALLMENTS = INSTALLMENTS_PURCHASES,
       PURCH_FREQ = PURCHASES_FREQUENCY,
       ONEOFF_FREQ = ONEOFF_PURCHASES_FREQUENCY,
       INSTALL_FREQ = PURCHASES_INSTALLMENTS_FREQUENCY,
       CASHADV_FREQ = CASH_ADVANCE_FREQUENCY,
       MIN_PAY = MINIMUM_PAYMENTS) %>%
```

```
# create categorical variable "AT_RISK", which will be used later as the response variable in the Cla
```

```
mutate(CREDIT_UTIL = round(BALANCE / CREDIT_LIMIT, 2),
       AT_RISK = ifelse(CREDIT_UTIL > 0.30, "TRUE", "FALSE"))
```

```
ccdata %>% filter(AT_RISK=="TRUE") %>%
```

```
summarize(customers_at_risk = n(), not_at_risk = 4475-customers_at_risk)
```

```
## # A tibble: 1 x 2
```

```
##   customers_at_risk not_at_risk
```

```
##           <int>         <dbl>
```

```
## 1             2347         2128
```

```
# check for NA values
```

```
ccdata %>% summarize_all(function(x)sum(is.na(x)))
```

```
## # A tibble: 1 x 20
```

```
##   CUST_ID BALANCE BAL_F~1 PURCH~2 ONEOFF INSTA~3 CASH_~4 PURCH~5 ONEOF~6 INSTA~7
```

```
##   <int>   <int>   <int>   <int>   <int>   <int>   <int>   <int>   <int>   <int>
```

```
## 1     0     0     0     0     0     0     0     0     0     0
```

```
## # ... with 10 more variables: CASHADV_FREQ <int>, CASH_ADVANCE_TRX <int>,
```

```
## #   PURCHASES_TRX <int>, CREDIT_LIMIT <int>, PAYMENTS <int>, MIN_PAY <int>,
```

```
## #   PRC_FULL_PAYMENT <int>, TENURE <int>, CREDIT_UTIL <int>, AT_RISK <int>, and
```

```
## #   abbreviated variable names 1: BAL_FREQ, 2: PURCHASES, 3: INSTALLMENTS,
```

```
## #   4: CASH_ADVANCE, 5: PURCH_FREQ, 6: ONEOFF_FREQ, 7: INSTALL_FREQ
```

```
## # i Use `colnames()` to see all variable names
```

```
# for variables "CREDIT_LIMIT" and "MINIMUM_PAY", replace NA values with mean
```

```
ccdata$CREDIT_LIMIT[is.na(ccdata$CREDIT_LIMIT)] <- mean(ccdata$CREDIT_LIMIT, na.rm = TRUE)
```

```
ccdata$MIN_PAY[is.na(ccdata$MIN_PAY)] <- mean(ccdata$MIN_PAY, na.rm = TRUE)
```

Cluster Analysis

```
library(cluster)
```

```
pam_dat <- ccdata %>% select(BAL_FREQ, PURCH_FREQ, INSTALL_FREQ, ONEOFF_FREQ, CASHADV_FREQ, MIN_PAY)
```

```
sil_width <- vector()
```


```
for(i in 2:10){
```

```
  pam_fit <- pam(pam_dat, k = i)
```

```

  sil_width[i] <- pam_fit$silinfo$avg.width
}
ggplot()+
  geom_line(aes(x=1:10,y=sil_width))+
  scale_x_continuous(name="k", breaks=1:10)

```



```

# looks like five clusters is best!

```

```

pam1 <- pam_dat %>% pam(k=5)
pam1$silinfo$avg.width # 0.6275

```

```

## [1] 0.6388471

```

```

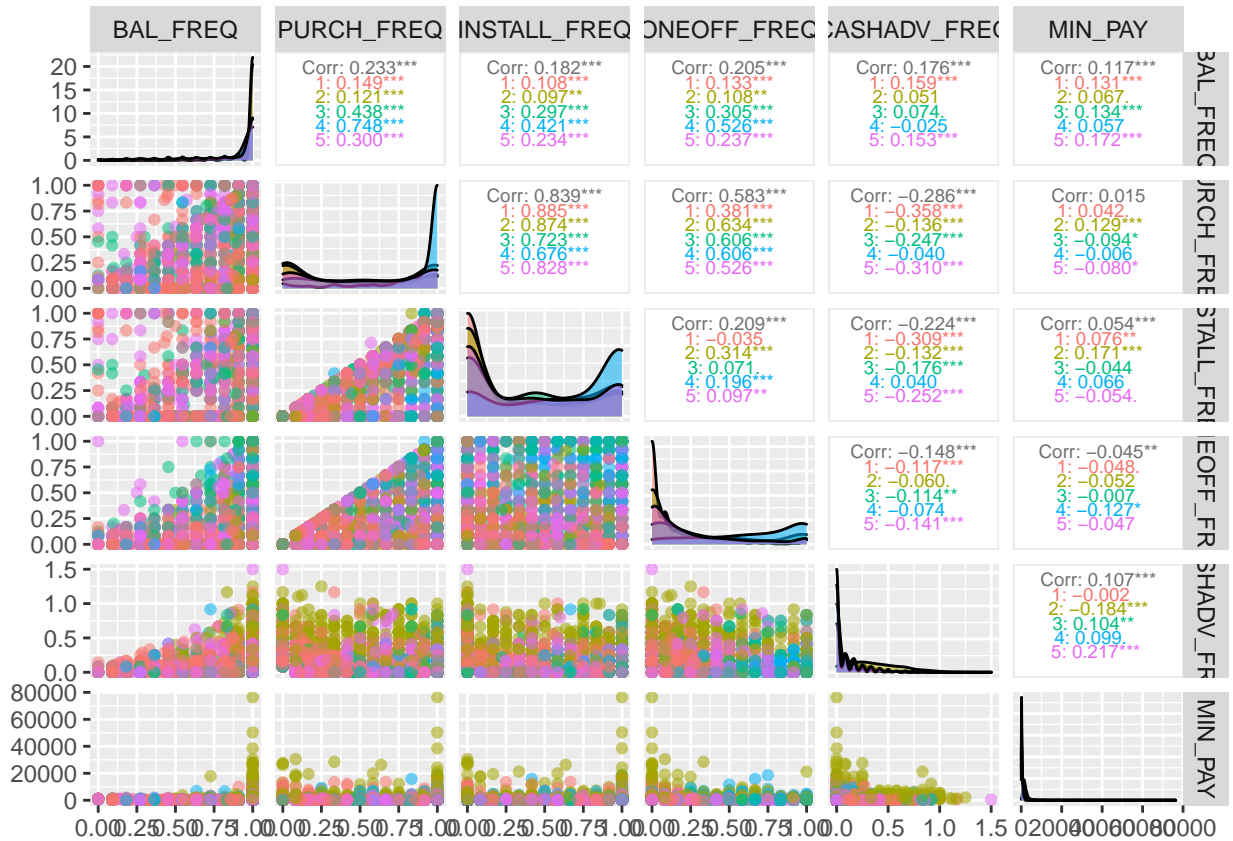
# cluster visualization

```

```

library(GGally)
clust <- ccddata %>% pam(k=5)
ccdata %>%
  select(BAL_FREQ, PURCH_FREQ, INSTALL_FREQ, ONEOFF_FREQ, CASHADV_FREQ, MIN_PAY) %>%
  mutate(cluster=as.factor(clust$clustering)) %>%
  ggpairs(columns=1:6, aes(color=cluster, alpha = 0.5),
    upper = list(continuous = wrap("cor", size = 2.5)))

```



I performed PAM clustering on the following variables: balance frequency, purchase frequency, installment frequency, oneoff frequency, cash advance frequency, and minimum payments. To determine the ideal number of k clusters, I calculated the silhouette width from k=2 to k=10. In my analysis, 5 clusters returned the highest average silhouette width, at 0.63. This value indicates a reasonable clustering structure.

Subsequently, I used ggpairs() to visualize all pairwise variable combinations and colored them by the 5 cluster assignments. The pair with the highest positive correlation is installment frequency and purchase frequency (corr=0.839), and the pair with the most negative correlation is cash advance frequency and purchase frequency (corr=-0.286).

Dimensionality Reduction with PCA

```
# scale data and summarize PCA results
ccdata %>% select(BAL_FREQ, PURCH_FREQ, INSTALL_FREQ, ONEOFF_FREQ, CASHADV_FREQ, MIN_PAY) %>% scale %>%
summary(ccdata_pca, loadings=T)

## Importance of components:
##               Comp.1   Comp.2   Comp.3   Comp.4   Comp.5
## Standard deviation  1.5211153 1.1266278 0.9775104 0.8701242 0.7971574
## Proportion of Variance 0.3857182 0.2115957 0.1592900 0.1262142 0.1059337
## Cumulative Proportion 0.3857182 0.5973138 0.7566039 0.8828181 0.9887517
##               Comp.6
## Standard deviation  0.25975852
## Proportion of Variance 0.01124826
## Cumulative Proportion 1.00000000
##
## Loadings:
```

```
##          Comp.1 Comp.2 Comp.3 Comp.4 Comp.5 Comp.6
## BAL_FREQ    0.224  0.620  0.330  0.124  0.664
## PURCH_FREQ    0.634                -0.191 -0.742
## INSTALL_FREQ  0.547        -0.271  0.512 -0.152  0.585
## ONEOFF_FREQ   0.431          0.414 -0.688 -0.252  0.325
## CASHADV_FREQ -0.249  0.602  0.271  0.256 -0.660
## MIN_PAY       0.502 -0.756 -0.420
```

```
# determine number of PCs to keep
```


```
eigval <- ccdata_pca$sdev^2
```

```
varprop=round(eigval/sum(eigval), 2)
```

```
eigval<-ccdata_pca$sdev^2
```

```
varprop=round(eigval/sum(eigval), 2)
```

```
ggplot() + geom_bar(aes(y=varprop, x=1:6), stat="identity") + xlab("") +
  geom_text(aes(x=1:6, y=varprop, label=round(varprop, 2)), vjust=1, col="white", size=5) +
  scale_y_continuous(breaks=seq(0, .5, .1), labels = scales::percent) +
  scale_x_continuous(breaks=1:6)
```




```
eigval # eigenvalues
```

```
##      Comp.1      Comp.2      Comp.3      Comp.4      Comp.5      Comp.6
## 2.31379178 1.26929031 0.95552661 0.75711613 0.63545990 0.06747449
```

```
# plot scores to show data with respect to 2 PCs
```

```
ccdataf <- data.frame(Customer=ccdata$CUST_ID, PC1=ccdata_pca$scores[,1], PC2=ccdata_pca$scores[,2])
```

```
ggplot(ccdataf, aes(PC1, PC2)) + geom_point()
```



```
library(factoextra)
fviz_pca_biplot(ccdata_pca)
```



```

# report a confusion matrix
table(actual = ccddata$AT_RISK, predicted = log_score > .5) %>% addmargins()

##          predicted
## actual  FALSE TRUE  Sum
##  FALSE  1792  336 2128
##   TRUE   244 2103 2347
##   Sum    2036 2439 4475

TNR <- 1792 / 2128 # 0.842 (Specificity)
FP <- 1 - TNR      # 0.158
TPR <- 2103 / 2347 # 0.896 (Sensitivity / Recall)
FN <- 1 - TPR      # 0.104

# Positive Predictive Value (PPV) / Precision
PPV <- TPR / (TPR + FP)
PPV #0.850

## [1] 0.8501851

# perform k-fold CV on this same model
set.seed(1234)
k=10
data <- ccddata[sample(nrow(ccddata)),] #randomly order rows
folds<-cut(seq(1:nrow(data)),breaks=k,labels=F) #create folds
diags<-NULL

for(i in 1:k){
  ## create training and test sets
  train<-data[folds!=i,]
  test<-data[folds==i,]
  truth<-test$AT_RISK

  ## train model on training set
  fit<-glm(AT_RISK == "TRUE" ~ ONEOFF + INSTALLMENTS + CASH_ADVANCE + BAL_FREQ + PURCH_FREQ + ONEOFF_FREQ,
    data=train, family="binomial")
  probs<-predict(fit,newdata = test,type="response")

  ## test model on test set (save all k results)
  diags<-rbind(diags,class_diag(probs,truth, positive="TRUE"))
}
summarize_all(diags,mean) # AUC: 0.9427

##          acc      sens      spec      ppv      f1      ba      auc
## 1 0.86996 0.89599 0.84104 0.86175 0.87833 0.86853 0.94275

```

For my classifier methods, I used “AT_RISK” as my response variable and the following variables as my predictors: ONEOFF, INSTALLMENTS, CASH_ADVANCE, BAL_FREQ, PURCH_FREQ, ONEOFF_FREQ, INSTALL_FREQ, CASHADV_FREQ, MIN_PAY, and PRC_FULL_PAYMENT. Logistic regression had an area under the curve (AUC) value of 0.9435, indicating that its predictive performance was very good. The confusion matrix also found that the logistic model could correctly classify AT_RISK (low credit score individuals) about 89.6% of the time, and non-risk individuals 84.2% of the time. Overall precision was about 85%. The AUC value for cross-validation (CV) was 0.9427; since this was only slightly lower than the original AUC value, the model did not appear to show signs of overfitting.

Non-Parametric Classifier

```
library(caret)

knn_fit <- knn3(AT_RISK == "TRUE" ~ ONEOFF + INSTALLMENTS + CASH_ADVANCE + BAL_FREQ + PURCH_FREQ + ONEOFF_F1)

y_hat_knn <- predict(knn_fit, ccddata)

class_diag(y_hat_knn[,2], ccddata$AT_RISK, positive="TRUE") # AUC: 0.9475

##          acc   sens   spec   ppv    f1    ba    auc
## Metrics 0.8708 0.8769 0.8642 0.8769 0.8769 0.8705 0.9475

# confusion matrix
table(actual = ccddata$AT_RISK, predicted = y_hat_knn[,2] > .5) %>% addmargins

##          predicted
## actual  FALSE TRUE  Sum
##  FALSE   1839  289 2128
##   TRUE    289 2058 2347
##    Sum     2128 2347 4475

TNR <- 1840 / 2128 # 0.865 (Specificity)
FP <- 1 - TNR      # 0.135
TPR <- 2058 / 2347 # 0.877 (Sensitivity / Recall)
FN <- 1 - TPR      # 0.123

# Positive Predictive Value (PPV) / Precision
PPV <- TPR / (TPR + FP)
PPV #0.866

## [1] 0.8662932

# k-fold CV on the same model
k=10
data<-ccdata[sample(nrow(ccdata)),]
folds<-cut(seq(1:nrow(ccdata)),breaks=k,labels=F)
diags<-NULL
for(i in 1:k){
  train<-data[folds!=i,]
  test<-data[folds==i,]
  truth<-test$AT_RISK

  fit<-knn3(AT_RISK == "TRUE" ~ ONEOFF + INSTALLMENTS + CASH_ADVANCE + BAL_FREQ + PURCH_FREQ + ONEOFF_F1)
  probs<-predict(fit,newdata = test)[,2]

  diags<-rbind(diags,class_diag(probs,truth, positive="TRUE"))
}
summarize_all(diags,mean) # AUC: 0.9012

##          acc   sens   spec   ppv    f1    ba    auc
## 1 0.84871 0.86918 0.82528 0.8462 0.85738 0.84724 0.90202
```

For my nonparametric model, I used k-nearest-neighbors (KNN). This method returned an AUC value of 0.9475, indicating good predictive performance. However, for its cross-validation, the AUC value was 0.9012, a noticeable decrease and potential result of overfitting. It appeared that the logistic method (linear classifier) had a better CV performance than KNN.

Regression/Numeric Prediction

```
# Fit a linear regression model or regression tree to your entire dataset, predicting one of your numeric variables
fit <- lm(AT_RISK == "TRUE" ~ ONEOFF + INSTALLMENTS + CASH_ADVANCE + BAL_FREQ + PURCH_FREQ + ONEOFF_FREQ)
probs <- predict(fit)
class_diag(probs, ccddata$AT_RISK, positive="TRUE") # AUC: 0.9035
```

```
##           acc   sens   spec   ppv   f1   ba   auc
## Metrics 0.8322 0.9433 0.7096 0.7818 0.855 0.8265 0.9036
```

```
lm_summary <- summary(fit)
```

```
# calculate MSE for the overall dataset
mean(lm_summary$residuals^2)
```

```
## [1] 0.138627
```

```
k=10
data<-ccdata[sample(nrow(ccdata)),]
folds<-cut(seq(1:nrow(ccdata)),breaks=k,labels=F)
diags<-NULL
for(i in 1:k){
  train<-data[folds!=i,]
  test<-data[folds==i,]
  truth<-test$AT_RISK

  fit<-lm(AT_RISK == "TRUE" ~ ONEOFF + INSTALLMENTS + CASH_ADVANCE + BAL_FREQ + PURCH_FREQ + ONEOFF_FREQ)
  probs<-predict(fit,newdata = test)

  diags<-rbind(diags,class_diag(probs,truth, positive="TRUE"))
}
summarize_all(diags,mean) # AUC: 0.9027
```

```
##           acc   sens   spec   ppv   f1   ba   auc
## 1 0.8286 0.94301 0.70182 0.77802 0.85241 0.82242 0.9025
```

The dataset overall had a mean standard error (MSE) of 0.1372. The linear regression had an AUC of 0.9035 and the cross-validation reported an AUC of 0.9027; this was only a slight decrease, so overfitting was unlikely, but overall this classifier method exhibited the weakest predictive performance.